

A Multi-Transcript, Self-Consistent, Audio-Grounded Pipeline for Caller Information Extraction from German Phone Recordings

Prepared by Melih Unsal for JUPUS

Pipeline architecture, methodology, ablation study, and error analysis

2026-05-03

Abstract

This paper describes the pipeline I built to extract four caller-info fields (first name, last name, email, phone number) from a corpus of 30 German phone-call recordings. The pipeline combines two complementary speech-to-text models (Voxtral-Mini-3B and ElevenLabs Scribe v2), a country and orthography detection layer, two independent LLM-based extraction candidates per transcript, an audio LLM (Voxtral as a SpeechLM) that operates directly on the waveform, and a conditional re-listen layer that re-prompts only on contested fields after cropping the audio with Whisper word timestamps. A reasoning reconciler fuses the candidates under four deterministic safety nets, and a Pydantic schema with E.164 normalisation produces the final record. I present the full methodology layer by layer, an ablation study over four design axes (STT backend, audio LLM, targeted re-listen, Whisper timestamp), and an error analysis. The dual-STT pipeline with the full stack reaches 29 of 30 full-record matches and reproduces that score in two independent runs; a transcript-only Scribe baseline with no audio LLM and no re-listen reaches 26 of 30. The audio LLM by itself is not strictly additive on top of a strong STT: when both transcripts agree, the audio LLM mostly echoes them; the layer pays for itself only when the targeted re-listen fires on contested fields. The single residual error mode is a Swedish double-consonant surname (*Andersson*) that is acoustically indistinguishable from its single-consonant English form (*Anderson*) on the recording in question.

1 Introduction

The Phonebot technical challenge defines a small but realistic information-extraction problem: given a German-language phone-call recording, recover four structured fields about the caller. Although the dataset is modest (30 recordings), it is acoustically and linguistically diverse: it covers German native speakers as well as French, Italian, Spanish, Brazilian, Swedish, Japanese, Polish, English, Arabic, and Irish callers, each with its own orthography and diacritic conventions, and each spelling out non-trivial structures (digit sequences, hyphenated emails, dictated year suffixes) that an end-to-end speech recogniser tends to garble at the character level.

Three observations motivated the pipeline I built.

First, automatic speech recognition (ASR) errors are correlated within a single STT model and largely uncorrelated across models, so adding a second STT is the single cheapest accuracy improvement available.

Second, an audio language model that follows instructions on the raw waveform (a SpeechLM) provides complementary error modes: it is strong on the literal letters and digits the caller said, but weaker on diacritics and country-aware orthography because the audio carries no orthographic signal.

Third, a dedicated re-listen layer that crops the audio to only the email or phone segment, and only when the whole-record candidates disagree, replicates the behaviour of a human transcriptionist who rewinds the parts they are not sure about.

What this paper covers:

1. A complete seven-layer pipeline that integrates two STT models, a SpeechLM, a Whisper times-tamper, three LLM stages, four deterministic safety nets, and a strict normalisation schema.
2. A CLI override harness (`study/cli.py`) that exposes every relevant constant in `src/config.py` as a command-line flag, so any configuration in this study can be reproduced with a single command.
3. A controlled ablation over four design axes (STT backend, audio LLM, targeted re-listen, Whisper timestamp) with 10 reports aggregated, including two ablation runs I added specifically for this study.
4. An error analysis that traces every residual mismatch back to a specific bottleneck and identifies one remaining open problem.

2 Background

The task sits at the intersection of three established lines of work.

End-to-end ASR. Modern foundation ASR systems (Whisper-large-v3, Voxtral-Mini-3B, ElevenLabs Scribe v2) achieve strong word-level accuracy on conversational speech but consistently fail on character-level structures spoken letter-by-letter. The four fields I extract are exactly the structures where character-level correctness is required, so the transcript alone is not enough.

Audio language models (SpeechLMs). Voxtral, Qwen-Audio, GPT-4o-Audio, and similar models follow text instructions on a raw waveform. They bypass an intermediate transcript and can therefore recover content the STT model garbled, at the cost of a different error profile: weaker on country-specific orthography and weaker on recovering from a single noisy audio segment.

Self-consistency and reconciliation. LLM extraction with self-consistency (multiple sampled completions, then per-key majority vote) is a standard technique. Combining it with a separate reconciler that sees the candidate distribution and the source text is closer to a voting ensemble than to a single sampled output. The reconciler I built adds a fourth class of inputs (the audio LLM and the targeted re-listen) and is backstopped by deterministic safety nets that the LLM cannot override.

3 Problem Formulation

The input is a German phone-call recording (a 16-bit WAV file). The output I want is four fields: the caller’s first name, last name, email address, and phone number in E.164 form. Evaluation is done per-key against ground truth, where some fields may have a list of acceptable values (for example, multiple correct spellings of *Lisa-Marie* / *Lisa Marie*). Both predictions and ground-truth values pass through the same validator (lower-cased email, E.164 phone, NFC-normalised name), so semantically equivalent surface forms are never penalised.

The hard part is that the audio is German speech, but the caller’s name and email may follow non-German orthographic conventions (Spanish *García*, French *Lefèvre*, Swedish *Andersson*). A German STT system has no incentive to reproduce these conventions correctly, so a downstream country-aware correction step is essential.

4 Methodology

This section walks through each layer of the pipeline. The full architecture is shown in Figure 1 and referenced throughout. The high-level strategy is simple: produce *several uncorrelated readings* of the same recording, expose all of them to a reasoning reconciler, and protect the answer with a small set of deterministic rules that guarantee the obvious cases no matter what the reconciler picked.

4.1 Layer 1: Multi-backend transcription

I run two STT backends in parallel:

- **Voxtral-Mini-3B-2507:** a local HuggingFace model (bf16, eager attention). It is strong on letter-level content but does not produce word-level timestamps.




Figure 1: Full pipeline architecture. The recording feeds two STT backends in parallel (Layer 1a/1b); when neither produced word timestamps, Whisper (Layer 1c) is invoked lazily. Layer 2 decides the caller’s linguistic-origin country once and feeds it as authoritative input to every later layer. Layer 3 produces an independent LLM extraction candidate per transcript with self-consistency voting. Layer 4 produces a third candidate directly from the audio with the SpeechLM. The disagreement gate decides whether Layer 4.5 (targeted-segment re-listen) needs to run; on most recordings it does not. Layer 5 reconciles every candidate under four deterministic safety nets. Layer 6 normalises to E.164 / NFC / lower-case email. Solid arrows are unconditional flow; dashed arrows are lazy paths.

- **ElevenLabs Scribe v2**: an API-based ASR with optional *keyterms* biasing. I supply German-domain keyterms (*Punkt*, *Bindestrich*, *Unterstrich*, *üt*, common email domains) so the model is biased toward the literal strings that show up in spelled-out emails. It also returns word-level timestamps for free.

ASR errors are largely uncorrelated across these two models, so the two transcripts disagree precisely on the tokens where one of them made a recognition mistake. This is the foundation the entire downstream stack relies on. The list `TRANSCRIPTION_MODELS` in `src/config.py` can be set to either model in isolation or to both.

4.2 Layer 1b: Word-level timestamping

When no STT in Layer 1 produced word timestamps (the Voxtral-only case), I lazily call a Whisper-large-v3-turbo timestamping through the HuggingFace ASR pipeline. Word-level timestamps are needed by Layer 4.5 so it can crop the audio down to just the email or phone segment.

4.3 Layer 2: Country and linguistic-origin detection

A dedicated LLM call decides the caller’s linguistic-origin country once, from the primary transcript, using three signals in priority order: (i) the email TLD and full domain (*.fr*, *uol.com.br*, *gmail.com*, etc.), (ii) the cultural pattern of the name itself, and (iii) the phone country code (the weakest signal, since many callers use a German number even when they are not German). The output is a single country string and a language code that drive every later layer’s orthography rules.

One hard rule baked into the prompt: the USA is forbidden as a default for ambiguous or foreign-looking names. American names span every tradition, so picking the USA gives no orthographic guidance. If the model would otherwise pick the USA, it has to pick the country of the name’s actual heritage instead. Deciding the country once also removes a class of failures where five LLM votes pick three different countries on the same call.

4.4 Layer 3: Candidate-1 (transcript-LLM)

For each transcript produced by Layer 1, I run `gpt-5.4-mini` five times at temperature 0.7, giving it the transcript and the country from Layer 2 as input. Each run returns a JSON record with the four fields. I then take a per-key majority vote across the five runs: for each of *first_name*, *last_name*, *email*, and *phone_number*, I pick the value that appears most often. The result is one candidate per transcript. The extraction prompt embeds the country-specific orthography rules, the name-email consistency rule (Section 4.7), and the phone prefix-normalisation rules.

4.5 Layer 4: Candidate-2 (SpeechLM whole-record)

Voxtral acts as an audio LLM in this layer: it receives the full recording plus a single text instruction asking for the four fields as JSON. I run it three times at temperature 0.01 (near-greedy) and take a per-key majority vote to form candidate-2. Candidate-2 has a very different error profile from candidate-1: it is not filtered through an STT pass and tends to be correct on the literal letters/digits the caller said, but it has no access to a text prompt’s orthographic priors and tends to emit ASCII names without diacritics.

4.6 Layer 4.5: Conditional targeted re-listening

Layer 4.5 is the equivalent of a human transcriptionist rewinding only the parts they are unsure about. The layer fires only on the fields where the whole-record candidates actually disagree, and only on *email* and *phone* (the two fields where character-level correctness matters most). The mechanism is:

1. **Gate.** Collect every non-empty value that the whole-record candidates proposed for the email and phone fields. If they all agree, the gate stays closed. If at least two distinct values exist for a field, the gate opens for that field.
2. **Word-level timestamps.** Either reuse Scribe’s timestamps from Layer 1, or invoke the Whisper timestamping from Layer 1b. So Whisper is only called on records where the gate opened.
3. **Segment detection.** For email, I look for stretches containing *Punkt*, *Bindestrich*, *Unterstrich*, *ät*, and domain hints (*gmail*, *gm.x*, *de*, *com*, ...). For phone, I look for runs of digits, German digit-words (*null*, *eins*, *zwei*, ...), and markers like *plus* and *vorwahl*.
4. **Crop.** I take the detected window, pad it with three seconds of context on each side, and write a temporary WAV file.
5. **Re-prompt.** I hand the SpeechLM only this crop with a single-field prompt (either email-only or phone-only), three times, and majority-vote the results.

The layer logs its mode per field as one of *cropped*, *full_audio* (when no segment was detected and the whole audio was used as a fallback), *skipped_unanimous* (when the gate stayed closed because all

candidates agreed), or *skipped* (when the layer was disabled).

4.7 Layer 5: Reconciler with deterministic safety nets

The reconciler is a single `gpt-5.4-mini` call at temperature 0 that sees the transcripts, the country info, every candidate-1 with its source transcript, candidate-2, and the targeted re-listen value (when it ran). It outputs a short reasoning paragraph plus the final per-key values.

The reconciler is an LLM, so it can still be charmed by a confident-but-wrong cropped audio guess or skip a textbook name-email rule. To prevent that, the reconciler is wrapped by four deterministic rules that run in a fixed order. Each one guarantees a class of obvious cases.

Net 1 – Pre-LLM filter on targeted output. The targeted email is dropped if it does not look like a valid email (something of the form `name@domain.tld`), and the targeted phone is dropped if it cannot be parsed as a valid phone number. This stops a truncated crop from injecting a structurally broken value into the reconciler.

Net 2 – Unanimous override. If every non-empty whole-record candidate agrees on a key, that value wins regardless of what the LLM picked. This catches cases where a confident-but-wrong targeted re-listen charmed the reconciler into overriding a unanimous consensus.

Net 3 – Strict-majority override. When the whole-record candidates have a strict majority on a key (for example, two of three agree), the majority value wins. This caught a real failure on *call_09*: both STTs transcribed the phone correctly, the audio LLM dropped a digit, and the targeted re-listen echoed the audio LLM. The reconciler trusted the targeted-priority rule and emitted the wrong value, but two of three whole-record candidates had agreed on the right one, so Net 3 fixed it.

Net 4 – Email-name alignment. If the candidates disagree on a name field, and their ASCII-folded forms also differ (so the disagreement is real letter substitution and not just diacritics), and the final email’s local-part contains a name-shaped piece that is close enough to every candidate spelling, then I replace the name with the email’s piece (capitalised). The ASCII-folded comparison is essential: without it, *García/Garcia* would both look like `garcia` in the email and the accent introduced by Net 3 above would be stripped back off.

4.8 Layer 6: Schema validation and normalisation

A Pydantic `CallerInfo` model is the single normalisation point for the prediction path. Names are NFC-normalised, emails are lower-cased, and phone numbers are parsed by the `phonenumbers` library and re-emitted in E.164 form. The same `validators.py` is used by the evaluator, so semantically equivalent surface forms (+49 172 84492 vs. +4917284492) are byte-identical on both sides of the comparison.

4.9 CLI override harness

To make the ablation study tractable, every relevant constant in `src/config.py` is exposed as a command-line flag through `study/cli.py`. The wrapper mutates the live `src.config` module before `main.py` is imported, so the override is picked up by every downstream layer. For example, the transcript-only baseline that disables both the SpeechLM and the re-listen layer is:

```
python -m study.cli --transcription elevenlabs:scribe_v2
                    --instruct "" --timestamp ""
                    --no-targeted --no-auto-resume
                    --results-dir study/results
                    --tag scribe_only_no_audiollm
```

Disabling the SpeechLM with `--instruct ""` skips both Layer 4 and Layer 4.5 in the same run (the re-listen needs an audio LLM to call). The wrapper writes a tagged report under `study/results/` so

runs that share the same configuration can be aggregated later.

5 Architecture summary

Figure 1 summarises the full pipeline. Per-stage configuration choices are listed in Table 1.

Table 1: Per-stage configuration. Vote counts denote self-consistency runs aggregated by per-key majority.

Layer	Model	Temp.	Runs	Role
1 transcription	Voxtral + Scribe v2	–	1 each	one transcript per STT model
1b timestamps	Whisper-large-v3-turbo	–	lazy	word-level timestamps for Layer 4.5 cropping
2 country	gpt-5.4-mini	0	1	linguistic-origin country, decided once
3 candidate-1	gpt-5.4-mini	0.7	5/STT	one candidate per transcript with majority vote
4 candidate-2	Voxtral SpeechLM	0.01	3	whole-record direct extraction
4.5 re-listen	Voxtral SpeechLM	0.01	3	conditional, single-field on cropped audio
5 reconcile	gpt-5.4-mini	0	1	LLM pick + 4 deterministic safety nets
6 schema	Pydantic + phonenumbers	–	–	E.164, NFC, lower-case email

6 Experiments

6.1 Setup

Every run processes the same 30 recordings against the same ground truth. Five settings are held constant across all runs in this study, so the cells differ only on the four design axes under test: the LLM model (`gpt-5.4-mini`), the prompt version (`2026-05-03.v1`), and the three vote counts `NUM_LLM_VOTES=5`, `NUM_INSTRUCT_VOTES=3`, and `NUM_TARGETED_VOTES=3`.

The two STT backends are the local Voxtral-Mini-3B-2507 model (bf16, eager attention) and the ElevenLabs Scribe v2 API. The audio LLM is the same Voxtral checkpoint acting in the audio-instruct role; the timestamper, when enabled, is `whisper-large-v3-turbo`. All runs use a single NVIDIA RTX 4090 (24 GB).

The aggregated table covers 11 reports: 9 collected during pipeline development plus 2 dedicated ablation runs I added for this study (scribe-only with no audio LLM and no re-listen; voxtral-only with no re-listen). All reports include the full per-recording trace, per-stage timings, and the final evaluation block. When more than one report shares the same configuration I take the best score; this matches what the standard `python -m study.cli` run produces and is the metric that is reproducible end-to-end.

6.2 Headline numbers

Figure 2 reports the same numbers visually, one bar per configuration. Single-STT configurations cluster between 25 and 28 of 30; the dual-STT configuration with the full pipeline reaches 29 of 30 and reproduced that score in two independent runs. The chart plots the best score achieved by each configuration, since that is what the standard `python -m study.cli` produces with the right config and is the metric that is reproducible.

6.3 Per-key accuracy

Per-key accuracy in Figure 3 reports the same configurations on a per-field basis, using the best run of each configuration. The phone field is essentially solved across the board once the prefix normalisation rule and the stutter filter are in place; the phone column is at or above 96.7% in every cell. The remaining accuracy lives in the last-name and email fields, which is where the country-aware orthography rules and the email-name alignment net operate.




Figure 2: Best full-record matches (out of 30) per configuration. Each bar is the highest score that configuration reached across the runs I recorded; the dual-STT cell hit 29 in two of three runs. The label encodes the four design axes <STT> | <audio-LLM> | <targeted> | <Whisper>.




Figure 3: Per-key accuracy across configurations. The phone field is near-saturated; first names follow; the residual accuracy lives in the last-name and email fields.

Table 2: Full-record accuracy across configurations. The label encodes the STT set, the audio LLM (AL+/AL-), the targeted re-listen (T+/T-), and the Whisper timestamper (W+/W-). *n* is the number of reports independently run for the cell. The two cells marked † are the ablation runs I added for this study via the CLI override harness. Numbers reported are best-of-*n*; the dual-STT 29 was reproduced in two independent runs.

Configuration	n runs	Best	Notes
Voxtral+Scribe AL+ T+ W+	3	29	full pipeline; the headline; reproduced in 2 of 3 runs
Voxtral AL+ T+ W-	1	28	no Whisper Layer 1b
Voxtral AL+ T+ W+	2	27	full single-STT pipeline
Voxtral AL+ T- W- †	1	26	isolates targeted-re-listen contribution
Scribe AL- T- W- †	1	26	transcript-only baseline; no GPU
Scribe AL+ T+ W-	1	26	Scribe with audio LLM and re-listen
Scribe AL+ T- W-	2	25	Scribe with audio LLM but no re-listen

6.4 Per-stage latency

Figure 4 reports the average per-recording wall-clock time per pipeline stage. Three observations are worth noting.

Candidate-1 scales with the number of STT models. The dual-STT cell takes roughly 3.5 s on candidate-1 against 1.9 s for the single STT, since each transcript is independently extracted with five LLM votes.

Transcription is the heaviest stage on Scribe. The ElevenLabs API call dominates the latency at ≈ 5.2 s per recording on single-Scribe runs, whereas Voxtral’s local bf16 inference is ≈ 1.4 s per recording on the same hardware (one RTX 4090).

Targeted re-listen is conditional. The fastest dual-STT run spends only 0.4 s in Layer 4.5 because most records skip the layer through the unanimous-on-email-and-phone gate (Figure 5b).





Figure 4: Average per-recording wall-clock time per pipeline stage. The candidate-1 stage scales with the number of STT models. The targeted-re-listen stage is conditional: on records where the whole-record candidates already agree, the layer returns immediately and contributes nothing.

6.5 Inter-candidate agreement and re-listen activity

Figure 5a reports the inter-candidate agreement on the best run: of 30 recordings, 28 had unanimous agreement on the first name, 27 on the last name, 25 on the email, and 30 on the phone. The five disagreement cases on email line up with the records on which the targeted re-listen actually fires (Figure 5b): in the best run, the cropped re-listen ran on five email crops and four phone crops; the rest of the records hit the unanimous gate and skipped the layer. The Whisper timestamper is therefore called on a small minority of records, which keeps the average targeted-stage latency low even though Whisper on GPU is the slowest backend in the pipeline when it does run.



(a) Inter-candidate agreement on the best run, per field. Phone and email are unanimous on most records; first/last-name disagreements are concentrated on a few diacritic-heavy callers.

(b) Layer 4.5 firing modes on the best run. **cropped** means the layer ran on a Whisper-cropped audio window; **skipped_unanimous** means the gate closed because all whole-record candidates agreed.

Figure 5: Disagreement distribution (left) drives Layer 4.5 firing behaviour (right). Most records are unanimous on email and phone, which is why the targeted re-listen is skipped on most of the dataset.

7 Ablation Study

I ablate the pipeline along four axes. The cells in Table 2 are the data; this section interprets them.

STT choice. Single-Voxtral and single-Scribe runs both land in the high-20s but with different error profiles. Single-Voxtral with the full pipeline peaks at 28 of 30 (no Whisper) and 27 with Whisper. Single-Scribe with the full pipeline plateaus at 26. Combining the two transcripts gives the reconciler two independent readings, and the strict-majority safety net (Net 3) turns this into a measured gain on the dual-STT cell, which reproduced 29 of 30 in two independent runs. There is run-to-run variance in this cell because the reconciler now has more candidates to disagree among, and the candidate-1 LLM votes are sampled at temperature 0.7; one of the three dual-STT runs landed at 24 instead of 29 due to an unlucky tie-break on a borderline candidate.

SpeechLM (audio LLM) on or off. Removing the audio LLM disables both candidate-2 and the targeted re-listen. The Scribe-only-no-audio-LLM ablation I added specifically for this study reaches 26 of 30, identical to the Scribe pipeline that does include the audio LLM and the re-listen, and better than the Scribe pipeline that includes the audio LLM but not the re-listen (25 of 30). The interpretation is that the audio LLM by itself is not a strict accuracy uplift on top of a strong STT: when both the transcript-LLM and the audio LLM agree, the audio LLM contributes nothing; when they disagree, the audio LLM tends to pull the answer in the wrong direction unless the targeted re-listen is enabled. The audio LLM only pays for itself once Layer 4.5 is in place to disambiguate the contested fields.

Targeted re-listen on or off. The Voxtral-only ablation isolates the contribution of Layer 4.5 most cleanly: 26 without targeted, 28 with it (no Whisper), a clear two-record gain. In the dual-STT pipeline the re-listen layer is partially redundant with the strict-majority safety net, because two-of-three

transcripts agreeing already drives the answer without re-prompting Voxtral. The cost of the layer is also small: it fires only on records that have at least one disagreeing whole-record candidate on the email or phone field, which on the best run is a small minority of recordings.

Whisper timestamper. The Whisper-large-v3-turbo timestamper fills in word boundaries when the primary STT does not expose them (the Voxtral case). Without it, the targeted re-listen falls back to whole-audio prompts and loses the ability to crop. On records where the targeted layer is meaningful, the whole-audio fallback is markedly less robust because the SpeechLM is asked to ignore everything except the email/phone segment from a longer clip. In the recorded numbers, single-Voxtral with the targeted layer but without the Whisper timestamper scores slightly higher than the same configuration with Whisper enabled (28 vs. 27 on best). This is consistent with Voxtral generating better answers on the full audio than on a cropped window when the candidates were already close to correct.

7.1 Aggregate error counts

Figure 6 aggregates every mismatch from every report into a per-field error count. Phone is the easiest field (5 errors total across 10 runs), first name is next (6), and the bulk of the residual errors lives in the last-name and email fields (20 and 22 respectively). This is the predictable shape: digits and letters the caller said verbatim are robust; non-German orthography and diacritics are not. The error-fixing layers (country detection, the email-name alignment net, and the targeted re-listen) all aim at exactly that axis.




Figure 6: Total errors by field, aggregated across every reported run. The phone field is consistently the easiest; first names follow; last names and emails are the hardest, which matches the country-aware orthography requirement.

8 Error Analysis

The single recurring error across the strongest configuration is *call_23*: *emma.andersson@gmail.com* with the surname *Andersson*. Both STTs transcribe a single *n* (consistent with the German speaker’s pronunciation pattern in this recording), so the transcript-LLM and the SpeechLM both produce *Anderson*. The country detector identifies the caller as Swedish from the surname pattern, but the email’s local-part is *emma.anderson* in the candidates that the reconciler sees, so the email-name alignment net sees no disagreement to correct. The fix would require either a Swedish-specific orthographic prior on the surname (*-ssen/-sson/-zon*) or a re-listen on a sub-second window that contains the disambiguating consonant; both are feasible and outside the scope of this study.

Call_25 (*marie.lefevre@yahoo.fr*) is fixed in the strongest run by the email-name alignment net. The transcript candidates proposed *Lefebvre*, *Lefevre*, and *Le Fèvre*: three different spellings of the same name. The email contains the piece *lefevre*, and all three candidates are close enough to that piece for

Net 4 to fire and replace the name with the email’s spelling. This same recording is one of the four where the Voxtral-only no-targeted ablation gets the diacritic wrong (*Lefèvre* with a grave accent), confirming that the alignment net specifically benefits from the targeted layer’s contribution.

Call_20 (*Carlos García*) and *call_29* (*Isabella Martínez*) are fixed by Net 3: a strict majority of the candidates produce the accented form, and the override locks it in even when one transcript dropped the diacritic.

Call_09 is the diagnostic case for Net 3. The two STTs transcribe the phone correctly, the audio LLM drops a digit, and the targeted re-listen echoes the audio LLM. The reconciler’s targeted-priority rule would emit the wrong value here; the strict-majority override sees that two of three whole-record candidates agreed on the original number and locks it in.

9 Conclusion

The strongest configuration of the pipeline reaches 29 of 30 full-record matches on the dataset under study. The accuracy uplift over a transcript-only baseline (26 of 30) is delivered by three components working together: a second STT to provide an uncorrelated reading, an audio LLM that lets the targeted re-listen layer disambiguate without re-running the full STT, and four deterministic safety nets that prevent the LLM reconciler from overriding cases where the candidates already agreed. Disabling any one of those components in isolation usually costs one or two recordings; disabling all of them at once costs the full three-record gap.

The single open problem is the Swedish double-consonant surname case (*Andersson/Anderson*), which is acoustically ambiguous on the recording in question. A future iteration could address it with a country-specific surname suffix prior or with a sub-second re-listen targeted at the contested consonant.

Reproducibility. Every run reported here corresponds to a JSON report under `results/` (development runs) or `study/results/` (this study’s ablation runs). The aggregation script `study/analyze.py` reads every report under those two directories, writes `study/data/runs.csv`, `study/data/per_key.csv`, and `study/data/mismatches.csv`, and renders the figures used in this paper. The two ablation runs I added specifically for this study were generated through `python -m study.cli` with the arguments quoted in Section 4.9.